

UNIVERSITY OF SYDNEY

MACHINE LEARNING AND DATA MINING

Classification Task

ASSIGNMENT 1


Authors:

TINJU ABRAHAM
CLAUDIO ARACENA
RAFAEL POSSAS

Lecturers:

FABIO RAMOS
ROMAN MARCHANT

May 9, 2016



THE UNIVERSITY OF
SYDNEY

Abstract

Categorization of large documents is one of the areas where Natural Language Processing can be used to create models that can help to predict different features of a text. The first step for applying such models would be to create a VSM (Vector Space Model) where the words of the given description could be represented as numbers. Hence, a predictive algorithm would be able to process the text as a matrix, usually with a very high number of dimensions, and predict the possible outcome according to the data provided. As the data could be considered of very high dimensions, some pre-processing and feature extractions techniques are needed to make the algorithm faster when running on a computer. This work comprises of trying to predict a category of a given application through its text description. The outcome can be very useful for processing text in application stores like Play Store and Apple store since the amount of data increases every day along with the needs to better categorize the applications in their correct genre. One of the goals was to evaluate the performance of different algorithms using a third party machine learning library and to implement the one that would perform the best in the given dataset. The logistic regression was the algorithm of choice as it had the best overall metrics / results across all tested algorithms. In depth analysis of the results and implementation of the algorithm will be provided in this work along with a brief discussion of future work.

1 Introduction

Supervised Learning is one of the main approaches of Machine Learning. Classification is an instance of supervised learning, where the goal is to learn a mapping from inputs x to outputs y , where $y \in \{1, 2, 3, \dots, C\}$, with C being the number of classes.

The aim of this classification task is to learn to classify unlabelled test data (predict the category) which consists of the features without class names/labels. The task should be solved by taking the training set which consists of large set of labelled training set and building a classifier from those instances. The Classifier would then able to classify the unlabelled examples based on the information learned from the training set.




Figure 1: Image Source: Raschka, S. (2014). Naive Bayes and Text Classification I-Introduction and Theory.

In machine learning, the observations are often known as instances, the explanatory variables are termed

features and the possible categories to be predicted are classes. The information related to the classification task explored in this report are given below.

Instances: A training set with 20,104 rows each corresponding to an app

Features: Training set has 13,626 columns each corresponding to tf-idf values which are extracted from words in the description of each app. If a word is found in the description of an app, it has a tf-idf value. On the other hand, its tf-idf value is zero if the word is not found in the description of the app.

Classes: The training set has labels associated with each row. There are 30 labels or classes associated with the apps. The test data should be labeled into these categories by the classifier.

2 Methods

2.1 Exploratory Data Analysis (EDA)

Before start any data manipulation, it is good practice and an effective way to extract insights about the data to do an exploratory data analysis. This process is about detecting and describing patterns, trends, and relations in data, motivated by certain purposes of investigation [1], often with visual methods.

In the case of this assignment a good question to answer is how the labels are distributed in all 30 categories. To extract this kind of information the use of a histogram is suitable. The figure 2 shows the histogram of the frequency of categories.




Figure 2: Frequency Histogram for classes in the dataset

It is possible to observe that there is a equitable distribution among almost all categories with a frequency range from 621 to 766 instances. However, there are four categories which their frequencies are lower than the rest with a range from 359 to 484.

2.2 Pre-processing

Machine learning algorithms learn from data. It is critical to feed them the right data for the problem it needs to solve. Even if the data provided is good, it is needed to make sure that it is in a useful scale, format and even meaningful features are included.

Dimensionality reduction is the process of converting a set of highly dimensional data into a set with lesser dimensions [2]. This reduction not only helps in the processing needs for running the algorithm but also reveals latent variables in the data. In most cases, applying these methods can help in solving machine learning problems as it obtains better features for classification or regression algorithms.

The set given for this task can be categorized as highly dimensional if we put every word in our vocabulary as a feature to be processed by our algorithm. There are 13,026 words in total, however, each application description only uses a few of those words, therefore, making our dataset sparse. For this work PCA (Principal Component Analysis) was used as a dimensionality reduction option. PCA is a technique for taking a dataset and reducing its dimensions by finding the direction in which the tuples line up best [3]. The first principal component is that where the variance is the higher as it can be illustrated in the Figure 3.




Figure 3: Principal Component Analysis (PCA) example

Looking at the data also helps to identify some other pre-processing techniques. For instance, the most frequent features (sum on column) were removed, as they could be seen as “stop-words” in the text and, therefore, would not contribute too much for predicting the document category. The dataset went from 13.626 features to 13.048 after applying the aforementioned pre-processing technique. The Figure 6 illustrates the pipeline of the entire pre-processing stage.




Figure 4: Dimensionality Reduction process

2.3 Learning Procedure

In order to prevent over fitting and better train the algorithm, the dataset was divided into disjoint training and cross-validation sets as illustrated in the figure 5




Figure 5: Learning Procedure

In the simplest case, testing sets are constructed by splitting the original dataset into more than one part. However, evaluations obtained in this case tend to reflect the particular way the data are divided up. Statistical sampling helps to get more accurate measurements, which is known by cross-validation [9].

The main goal of cross-validation is to ensure that every sampling from the original set has the probability of appearing in the training and testing set. The k-fold cross validation strategy is that where the original sample is randomly partitioned into k subsamples and one is left out each iteration [9]. This method was chosen as in our tests it helped to avoid over fitting and better train the algorithm.

2.4 Classifier

The term Classifier refers to the classification algorithm that maps an unlabeled instance into a category or class. The classifier's evaluation is most often based on the accuracy of the prediction. The algorithm implemented for classification of the test data is explained below. For this work, Logistic Regression was implemented as it gave us the best accuracy when running the dataset with a 10-fold cross validation.

2.4.1 Logistic Regression

Logistic regression sometimes called the logistic model or logit model, analyzes the relationship between multiple independent variables and a categorical dependent variable, and estimates the probability of occurrence of an event by fitting data to a logistic curve [4]. It is based on the below discriminative model for binary classification.

$$p(y|\mathbf{x}, \mathbf{w}) = \text{Ber}(y|\sigma(\eta)) = \sigma(\eta)^y(1 - \sigma(\eta))^{1-y} \quad (1)$$

$$\eta = \mathbf{w}^T \mathbf{x} \quad (2)$$

$\sigma(\eta)$ refers to the sigmoid function, also known as the logistic or logit function. The term “sigmoid” means S-shaped. It is also known as a squashing function, since it maps the whole real line to $[0,1]$, which is necessary for the output to be interpreted as a probability. This is defined as

$$\sigma(\eta) = \frac{1}{1 + \exp(-\eta)} = \frac{e^\eta}{e^\eta + 1} \quad (3)$$

Putting these two steps together we get $p(y|\mathbf{x}, \mathbf{w}) = \text{Ber}(y|\sigma(\mathbf{w}^T \mathbf{x}))$. This is called logistic regression due to its similarity to linear regression [5]. In logistic regression training, weights are set to maximize the Conditional Log Likelihood (CLL).

$$\mathbf{W} \leftarrow \underset{\mathbf{W}}{\operatorname{argmax}} \sum_{d \in D} \ln P(Y^d | X^d, \mathbf{W}) \quad (4)$$

To prevent overfitting, regularization can be used by penalizing large weights by changing the training objective:

$$\mathbf{W} \leftarrow \underset{\mathbf{W}}{\operatorname{argmax}} \sum_{d \in D} \ln P(Y^d | X^d, \mathbf{W}) - \frac{\lambda}{2} \|\mathbf{W}\|^2 \quad (5)$$

Where λ is a constant that determines the amount of smoothing.

2.4.2 Implementation

Listing 1: Class Definition

```

1 import numpy as np
2 from scipy.optimize import fmin_l_bfgs_b
3 from joblib import Parallel, delayed
4 import multiprocessing
5
6 class LogisticRegression:
7     def __init__(self):
8         print("Logistic Regression Class created")
9         self.all_theta = []
10        self.list_classes = []

```

Listing 2: Sigmoid Function

```

1 def sigmoid(self, X):
2     return 1 / (1 + np.exp(-X))

```

Listing 3: Cost Function

```
1 def cost_function_reg(self, theta, X, y, l):
2     m, n = X.shape
3     J = (1/m) * (-y.T.dot(np.log(self.sigmoid(X.dot(theta)))) \
4         - (1-y.T).dot(np.log(1 - self.sigmoid(X.dot(theta))))) \
5         + (1/m)* 0.5 * theta[1:].T.dot(theta[1:])
6     return J
```

Listing 4: Gradient Function

```
1 def grad_function_reg(self, theta, X, y, l):
2     m, n = X.shape
3     grad = (1/m) * X.T.dot(self.sigmoid(X.dot(theta)) - y)
4     grad[1:] = grad[1:] + (1/m)*theta[1:]
5     return grad
```

Listing 5: Training Process

```
1 def fit(self, X, y, l):
2     self.list_classes = list(set(y))
3     self.list_classes.sort()
4
5     classes = len(self.list_classes)
6     X = self.add_theta0(X)
7     num_cores = multiprocessing.cpu_count() -1
8     results = Parallel(n_jobs=num_cores)
9         (delayed(self.logistic_train_one_class)
10          (X, y, self.list_classes, l, c) for c in range(classes))
11     self.all_theta = np.asarray(results)
12
13 def logistic_train_one_class(self, X, y, list_classes, l, c):
14     m, n = X.shape
15     initial_theta = np.zeros(n)
16     y_class = self.get_y_class(y, list_classes, c)
17
18     def decorated_cost(theta):
19         return self.cost_function_reg(theta, X, y_class, l)
20
21     def decorated_grad(theta):
22         return self.grad_function_reg(theta, X, y_class, l)
23
24     theta = fmin_l_bfgs_b(decorated_cost, initial_theta, maxiter=50,
25                          fprime=decorated_grad)
26     return theta[0]
```

```

1  def predict(self, X):
2      m, n = X.shape
3      X = self.add_theta0(X)
4      y_pred = []
5      for i in range(m):
6          max_index = np.argmax(self.sigmoid(
7              self.all_theta.dot(np.transpose(X[i, :]))))
8          y_pred.append(self.list_classes[max_index])
9      return y_pred

```

3 Experiments and results

3.1 Metrics

Evaluating performance of a classifier can be as easy as measuring its accuracy, however, this method can lead to misleading conclusions if used as the only driver. Accuracy is not the only metric for evaluating the effectiveness of a classifier. In this work, it was also introduced the precision, recall and f-score metrics as they tell different things about the data.

Precision is the number of True Positives divided by the number of True Positives and False Positives, it is also called the Positive Predictive Value [6]. Recall is the number of True Positives divided by the number of True Positives and the number of false Negatives, it is also called Sensitivity or the True Positive Rate. Finally, the F1 Score conveys the balance between the precision and the recall.

$$\text{Precision } P = \frac{tp}{tp + fp} \quad (6)$$

$$\text{Recall } R = \frac{tp}{tp + fn} \quad (7)$$

$$\text{Accuracy } Acc = \frac{tp + tn}{tp + tn + fp + fn} \quad (8)$$

$$\text{F - score } F_1 = \frac{2PR}{P + R} \quad (9)$$

3.2 Experiments

All the experiments were done using the Cross Validation with 10 folds and the metrics precision, recall, accuracy, F1 Score and execution time were extracted as the tables 3.2 and 3.2 show. The tests were divided into the set with minimal dimensionality reduction (only stop words) and a set with full dimensionality reduction (pca + stop words). Even though dimensionality reduction looks like the way to go for this kind of problem, it did not bring better results or faster execution times in most cases.

Algorithm	Precision / STD (%)	Recall / STD (%)	F-Score / STD (%)	Accuracy (%)	Time (s)
Logistic Regression	65.62 \pm 1.28	66.37 \pm 1.30	65.37 \pm 1.28	65.34	1953
Gaussian Naïve Bayes	44.33 \pm 1.11	45.09 \pm 1.25	44.29 \pm 1.13	44.32	99
Multinomial Naïve Bayes	61.94 \pm 0.72	65.70 \pm 0.93	61.14 \pm 0.75	60.11	17
KNN	33.57 \pm 1.67	55.46 \pm 2.47	38.07 \pm 1.95	45.23	7091

Table 1: Metrics for 10 fold Cross Validation without PCA

Algorithm	Precision / STD (%)	Recall / STD (%)	F-Score / STD (%)	Accuracy (%)	Time (s)
Logistic Regression	61.41 \pm 0.75	60.88 \pm 0.75	59.98 \pm 0.74	61.69	2873
Gaussian Naïve Bayes	14.90 \pm 1.07	31.35 \pm 1.72	15.38 \pm 1.04	22.34	646
Multinomial Naïve Bayes	NA	NA	NA	NA	NA
KNN	39.68 \pm 1.03	56.00 \pm 1.34	43.57 \pm 0.82	44.38	2012

Table 2: Metrics for 10 fold Cross Validation with PCA

3.3 Confusion Matrix Sample




Figure 6: Confusion Matrix Sample

4 Discussion

While doing the classification task, the strategy was to do a comparison study of performance and accuracy using the existing packages of classification and then implement the objective function. In general, the following challenges were faced during the entire process.

Curse of dimensionality – When the number of dimension grows, the volume of the space increases so fast that the available data becomes sparse. The complexity of many classification algorithms is exponential with respect to the number of dimensions. Some of the algorithms tried in this experiment performed really bad because of this dimensionality issue. Few algorithms which we had to drop because of poor performance are listed below Xgboost – Poor performance that we had to drop experimenting with the algorithm. GradientBoostClassifier Random forest

Dimensionality Reduction

Combining different algorithms - After a better understanding of the strengths and limitations of each method, we wanted to investigate the possibility of integrating two or more algorithms in the classifier to utilize the strengths of one method to complement the weaknesses of another. Considering the huge dataset and the computational performance, it was difficult to do more experiments to come up with a feasible solution within the limited time of the assignment task.

5 Conclusions and future work

Few tasks that were not feasible in the limited time for the classification assignment, which we would like to try in future:

Would combining unrelated algorithms which were well performing and giving accurate results yield a better result?

Which feature selection methods are high performing and scalable across different classifiers in the mentioned classification task?

References

- [1] Natalia Andrienko and Gennady Andrienko. *Exploratory analysis of spatial and temporal data: a systematic approach*. Springer Science & Business Media, 2006.
- [2] Stuart Russell, Peter Norvig, and Artificial Intelligence. A modern approach. *Artificial Intelligence*. Prentice-Hall, Englewood Cliffs, 25:27, 1995.
- [3] Anand Rajaraman, Jeffrey D Ullman, Jeffrey David Ullman, and Jeffrey David Ullman. *Mining of massive datasets*, volume 1. Cambridge University Press Cambridge, 2012.
- [4] Hyeoun Park. An introduction to logistic regression: from basic concepts to interpretation with particular attention to nursing domain. *Journal of Korean Academy of Nursing*, 43(2):154–164, 2013.
- [5] Kevin P Murphy. *Machine learning: a probabilistic perspective*. MIT press, 2012.
- [6] Ron Kohavi et al. A study of cross-validation and bootstrap for accuracy estimation and model selection. In *Ijcai*, volume 14, pages 1137–1145, 1995.
- [7] Wikipedia. F1 score — wikipedia, the free encyclopedia, 2016. [Online; accessed 9-May-2016].
- [8] Wikipedia. Precision and recall — wikipedia, the free encyclopedia, 2016. [Online; accessed 9-May-2016].
- [9] Wikipedia. Statistical classification — wikipedia, the free encyclopedia, 2016. [Online; accessed 9-May-2016].
- [10] Wikipedia. Curse of dimensionality — wikipedia, the free encyclopedia, 2016. [Online; accessed 9-May-2016].